

Lecture 1.2

Reinforcement Learning: Dynamic Programming Methods

Alexey Gruzdev
alexey.s.gruzdev@gmail.com
HSE, Winter 2019

Recap: Markov Decision Process

- A Markov decision process (MDP) is a Markov reward process with decisions. It is an *environment* in which all states are Markov.
- A Markov reward process is a Markov chain with values.
- A *Markov Decision Process* is a tuple (S, A, P, R, γ)
 - S is a (finite) set of states
 - A is a finite set of actions
 - $P^a_{ss'}$ is a state transition probability matrix
 - $P^a_{ss'} = P[S_{t+1} = s' \mid S_t = s, A_t = a]$
 - R is a reward function, $R^a_s = E[R_{t+1} \mid S_t = s, A_t = a]$
 - γ is a discount factor, $\gamma \in [0, 1]$

Ergodic MDPs

Definition:

An MDP is ergodic if the Markov Chain induced by any policy is ergodic.

For any policy π , an ergodic MDP has an *average reward per time-step* ρ^π that is *independent* of start state.

$$\rho^\pi = \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E} \left[\sum_{t=1}^T R_t \right]$$

Recap: Tabular Cross Entropy method


- Policy is matrix A
 - $\pi(a \mid s) = P[A_t = a \mid S_t = s] = A_{s,a}$
- Sample N sessions with that policy
- Get M best sessions (elites)
- Elite = $[(s_1, a_1), (s_2, a_2), \dots, (s_k, a_k)]$
- Update policy:

$$\pi(a \mid s) = \frac{\text{took } a \text{ at } s \text{ state}}{\text{was at } s \text{ state}} = \frac{\sum [s_t = s][a_t = a]}{\sum [s_t = s]}$$

Recap: Cross Entropy Method

■ CEM:

- Evolutionary
- Go in the direction where elite goes.
- Easy to implement
- Black Box
- Need to play full episode to start learning



What is Dynamic Programming?

- *Dynamic* sequential or temporal component to the problem
- *Programming* optimizing a `program`, i.e. a policy
 - A method for solving complex problems
 - By breaking them down into sub problems
 - Solve the sub problems
 - Combine solutions to sub problems

Requirements for Dynamic Programming

- Dynamic Programming is a very general solution method for problems which have two properties:
- Optimal substructure
 - *Principle of optimality* applies
 - Optimal solution can be decomposed into sub problems
- Overlapping sub problems
 - Sub problems recur many times
 - Solutions can be cached and reused
 - Markov decision processes satisfy both properties
 - Bellman equation gives recursive decomposition
 - Value function stores and reuses solutions

Planning by Dynamic Programming

- Dynamic programming assumes full knowledge of the MDP
- It is used for *planning* in an MDP
- For prediction:
 - *Input*: **MDP** (S, A, P, R, γ) and policy π
 - or: **MRP** (S, P, R, γ)
 - *Output*: value function v_π
- For control:
 - *Input*: **MDP** - (S, A, P, R, γ)
 - *Output*: optimal value function v_* and optimal policy π_*

A bunch of applications for Dynamic Programming

- Dynamic programming is used to solve many other problems:
 - Scheduling algorithms
 - String algorithms (e.g. sequence alignment)
 - Graph algorithms (e.g. shortest path algorithms)
 - Graphical models (e.g. Viterbi algorithm)
 - Bioinformatics (e.g. lattice models)
- Q: Which algorithms do you know for finding shortest path problems in graphs?

State-value function

- The *return* G_t is the total discounted reward from time-step t .

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- The *state-value function* $v_{\pi}(s)$ of an MDP is the expected return starting from state s , and then following policy π

$$\begin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi}[G_t \mid S_t = s] = \mathbb{E}_{\pi}[R_t + \gamma G_{t+1} \mid S_t = s] = \\ &= \sum_a \pi(a \mid s) \sum_{r,s'} p(r, s' \mid s, a) [r + \gamma \mathbb{E}_{\pi}[G_{t+1} \mid S_{t+1} = s']] \\ &= \sum_a \pi(a \mid s) \sum_{r,s'} p(r, s' \mid s, a) [r + \gamma v_{\pi}(s')] \end{aligned}$$

$\pi(a \mid s)$ – policy stochasticity

$p(r, s' \mid s, a)$ – environment stochasticity


- Intuition: value of following policy π from state s

Bellman Expectation Equation #1

- Recursive definition of $v_{\pi}(s)$ is an important concept in RL

$$v_{\pi}(s) = \sum_a \pi(a | s) \sum_{r, s'} p(r, s' | s, a) [r + \gamma v_{\pi}(s')]$$

$$= E_{\pi} [R_t + \gamma v_{\pi}(s_{t+1}) | S_t = s]$$



Action-value function

- The *action-value function* $q_{\pi}(s, a)$ is the expected return starting from state s , taking action a , and then following policy π

$$\begin{aligned} q_{\pi}(s, a) &= \mathbb{E}_{\pi} [G_t \mid S_t = s, A_t = a] \\ &= \mathbb{E}_{\pi} [R_t + \gamma G_{t+1} \mid S_t = s, A_t = a] \\ &= \sum_{r, s'} p(r, s' \mid s, a) [r + \gamma \mathbb{E}_{\pi} [G_{t+1} \mid S_{t+1} = s']] \\ &= \sum_{r, s'} p(r, s' \mid s, a) [r + \gamma v_{\pi}(s')] \end{aligned}$$

- Intuition: value of following policy π from state s , after committing action a

Relations between Action/state value functions

- We know relationship $q_{\pi}(s, a)$ in terms of $v_{\pi}(s)$

$$q_{\pi}(s, a) = \sum_{r, s'} p(r, s' | s, a) [r + \gamma v_{\pi}(s')]$$

- What about $v_{\pi}(s)$ in terms of $q_{\pi}(s, a)$

$$v_{\pi}(s) = \sum_a \pi(a | s) \sum_{r, s'} p(r, s' | s, a) [r + \gamma v_{\pi}(s')]$$


$$v_{\pi}(s) = \sum_a \pi(a | s) q_{\pi}(s, a)$$

Bellman Expectation Equation #2

- Recursive definition of $q_\pi(s, a)$ in term of $q_\pi(s, a)$

$$q_\pi(s, a) = \sum_{r, s'} p(r, s' | s, a) [r + \gamma \sum_{a'} \pi(a' | s') q_\pi(s', a')]]$$

Backup
diagram
for $q(s, a)$



Recap: optimal policy

- Define a partial ordering over policies:

$$\pi \geq \pi' \text{ if } v_\pi(s) \geq v_{\pi'}(s) \quad \forall s$$

Theorem: *For any Markov Decision Process*


- *Best policy π_* performs better or equal than any other policy $\pi_* \geq \pi \quad \forall \pi$*
- *Optimal state-value function:*

$$v_*(s) = \max_{\pi} (v_{\pi}(s))$$


- *Optimal action-value function:*

$$q_{\pi_*}(s, a) = \max_{\pi} (q_{\pi}(s, a))$$

Bellman Optimality Equation #1



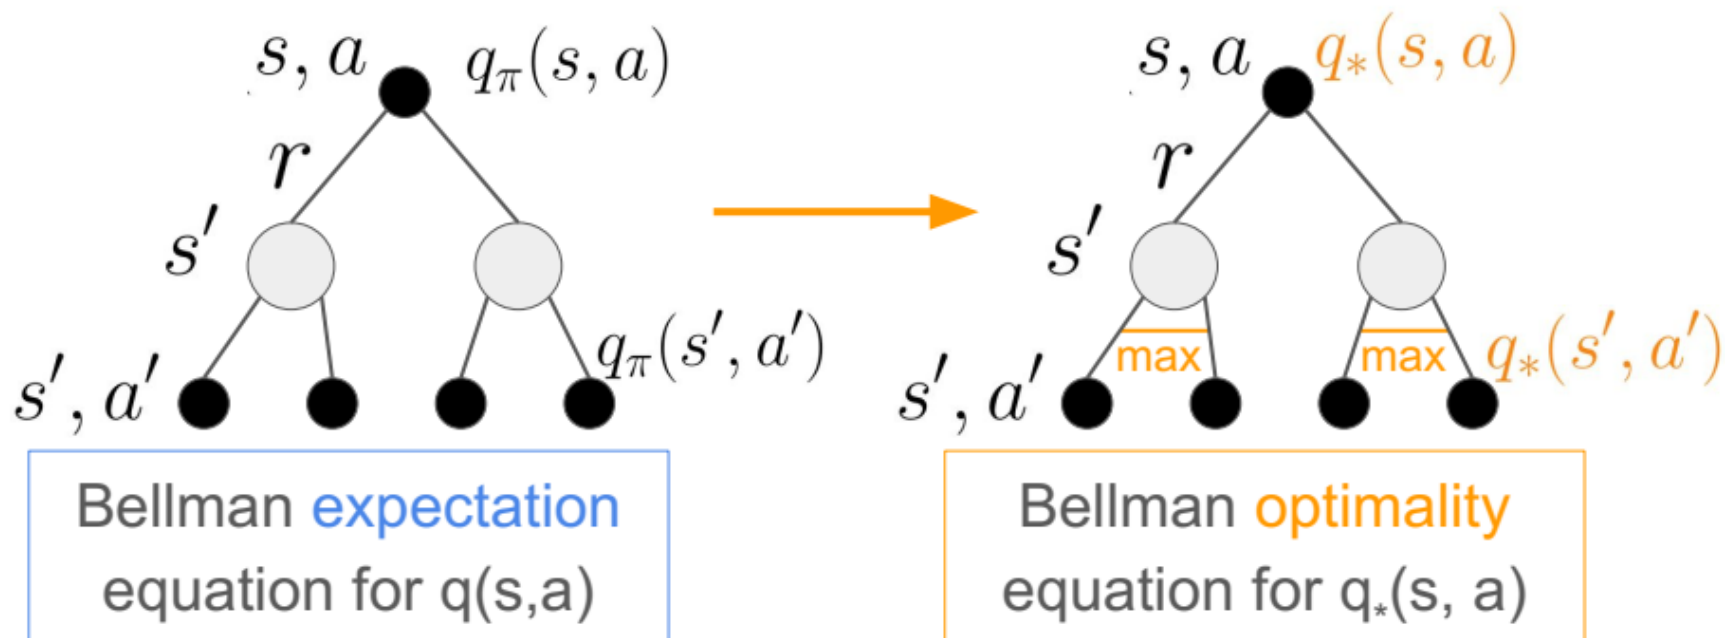
Bellman **expectation**
equation for $v(s)$



Bellman **optimality**
equation for $v_*(s)$

$$v_*(s) = \max_a \sum_{r,s'} p(r, s' | s, a) [r + \gamma v_*(s')]$$

Bellman Optimality Equation #2



$$q_*(s, a) = \sum_{r, s'} p(r, s' | s, a) [r + \gamma \max_{a'} q_*(s', a')]]$$

Iterative Policy Evaluation

- Problem: evaluate a given policy π
- Solution: iterative application of Bellman expectation backup
- $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_\pi$
- Using *synchronous* backups:
 - At each iteration $t + 1$
 - For all states $s \in S$
 - Update $V_{t+1}(s)$ from $V_t(s')$, where s' is a successor state of s
- Convergence is guaranteed but is not subject of this lecture.

Policy Evaluation

Iterative Policy Evaluation, for estimating $V \approx v_\pi$

Input π , the policy to be evaluated

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

Policy Improvement

- Given a policy π :

- Evaluate the policy π


$$v_{\pi}(s) = E [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t=s]$$

- Improve the policy by acting greedily with respect to v_{π}

$$\pi' = \text{greedy}(v_{\pi})$$


- In general, need more iterations of improvement / evaluation
- But this process of policy iteration always converges to π^*

Policy Improvement



Policy evaluation Estimate v_π
 Iterative policy evaluation

Policy improvement Generate $\pi' \geq \pi$
 Greedy policy improvement



Policy Iteration

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s', r | s, \pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)

3. Policy Improvement

policy-stable \leftarrow true

For each $s \in \mathcal{S}$:

old-action $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$

If *old-action* $\neq \pi(s)$, then *policy-stable* \leftarrow false

If *policy-stable*, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Modified Policy Iteration

- Does policy evaluation need to converge to v_π ?
- Or should we introduce a stopping condition
 - e.g. ε -convergence of value function
- Or simply stop after k iterations of iterative policy evaluation?
- Why not update policy every iteration? i.e. stop after $k = 1$
 - This is equivalent to *value iteration* (next section)

Value Iteration

- Problem: find optimal policy π
- Solution iterative application of Bellman optimality backup
- $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$
- Using synchronous backups
 - At each iteration $t + 1$
 - For all states $s \in S$
 - Update $v_{t+1}(s)$ from $v_t(s')$, where s' is a successor state of s
 - Unlike policy iteration, there is no explicit policy
 - Intermediate value functions may not correspond to any policy

Value Iteration

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

| $\Delta \leftarrow 0$

| Loop for each $s \in \mathcal{S}$:

| $v \leftarrow V(s)$

| $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$

| $\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$

Output a deterministic policy, $\pi \approx \pi_*$, such that

$$\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

Value vs Policy Iteration

- VI is faster per iteration – $O(|A||S|^2)$
- VI requires many iterations
- PI is slower per iteration – $O(|A||S|^2 + |S|^3)$
- PI requires few iterations
- No silver bullet → experiment with # of steps spent in policy evaluation phase to find the best

Resume

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

Let's do something real?!

Questions?

