

Lecture 1.3

Reinforcement Learning: Model-Free RL methods

Alexey Gruzdev
alexey.s.gruzdev@gmail.com
HSE, Winter 2019

Recap: Dynamic Programming


- $v_{\pi}(s), v_*(s)$
- If you know $v_*(s), p(r, s' | s, a) \rightarrow$ know optimal policy
- We can learn $v_*(s)$ with Dynamic Programming:

$$v_*(s) = \max_a \sum_{r, s'} p(r, s' | s, a) [r + \gamma v_*(s')]$$



- $q_{\pi}(s, a), q_*(s, a)$

$$q_*(s, a) = \sum_{r, s'} p(r, s' | s, a) [r + \gamma \max_{a'} q_*(s', a')]]$$

Decision making: reality check



Decision making: reality check



Model-Free Setup

- We don't know internal environment representation, e.g.

$$p(r, s' \mid s, a) - \text{unknown}$$


What should we do?

Let's count all letters that we introduced

- $a, r, s, p(r, s' | s)$
- $G_t(s)$
- $v_{\pi}(s), v_*(s)$
- $q_{\pi}(s, a), q_*(s, a)$
- π, π^*
- *Not enough? Next lectures we'll fix that!*

Learning from trajectories

- $s_1 \rightarrow a_1 \rightarrow r_1 \rightarrow s_2 \rightarrow \dots \rightarrow s_n$ – trajectory




- Model-based setup:
 - you can apply Dynamic Programming
 - you can plan (!)
- Model-free setup:
 - you can experiment with different actions
 - no guaranties (!!!)

Learning from trajectories

- $s_1 \rightarrow a_1 \rightarrow r_1 \rightarrow s_2 \rightarrow \dots \rightarrow s_n$ – trajectory
- *We can sample trajectories (a lot of trajectories!)*
- *What should we learn ?*
 - $p(r, s' \mid s, a)$
 - $v_\pi(s)$
 - $q_\pi(s, a)$

Monte-Carlo RL

- Just like N+1 heuristic:
 - Get all trajectories containing particular (s, a)
 - Estimate $G_t(s, a)$ for each trajectory
 - Average them to get *estimation* of expectation



Monte-Carlo RL

- MC methods learn directly from episodes of experience
- MC is *model-free*: no knowledge of MDP transitions / rewards
- MC learns from *complete* episodes: no bootstrapping
- MC uses the simplest possible idea: value = mean return
- Note: can only apply MC to *episodic* MDPs
 - All episodes must terminate

Incremental Mean

- The mean $\mu_1, \mu_2, \dots, \mu_k$ of a sequence x_1, x_2, \dots, x_k can be computed incrementally:

$$\begin{aligned}\mu_k &= \frac{1}{k} \sum_{j=1}^k x_j \\ &= \frac{1}{k} \left(x_k + \sum_{j=1}^{k-1} x_j \right) \\ &= \frac{1}{k} (x_k + (k-1)\mu_{k-1}) \\ &= \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})\end{aligned}$$

Temporal Difference Learning

- TD methods learn directly from episodes of experience
- TD is *model-free*: no knowledge of MDP transitions / rewards
- TD learns from *incomplete* episodes, by *bootstrapping*
- TD updates a guess towards a guess

Temporal Difference

- Just like in the 'incremental mean' example we can improve $q_{\pi}(s, a)$ iteratively:

$$q_*(s, a) = \sum_{r, s'} p(r, s' | s, a) [r + \gamma \max_{a'} q_*(s', a')]]$$

- We don't have $p(r, s' | s, a)$ to compute 'fair' expectation, so what should we do?


Temporal Difference

$$\begin{aligned} \sum_{r,s'} p(r, s' \mid s, a) [r + \gamma \max_{a'} q_*(s', a')] &\approx \\ \approx \frac{1}{N} \sum_i r_i + \gamma \max_{a'} Q(s'_i, a') \end{aligned}$$

- One more trick: use incremental averaging with 1 sample.


$$Q(s_t, a_t) = \alpha (r_t + \gamma \max_{a'} Q(s_{t+1}, a')) + (1 - \alpha) Q(s_t, a_t)$$

Q-learning




- Works on a sequence of
 - states (s)
 - actions (a)
 - rewards (r)

Q-learning




- Initialize $Q(s, a)$ with zeros
- Cycle:
 - Sample $\langle s, a, r, s' \rangle$ from environment

Q-learning



- Initialize $Q(s, a)$ with zeros
- Cycle:
 - Sample $\langle s, a, r, s' \rangle$ from environment
 - Compute $\hat{Q}(s, a) = r(s, a) + \gamma \max_{a_i} Q(s', a_i)$
 - Update: $Q(s_t, a_t) = \alpha \hat{Q}(s, a) + (1 - \alpha) Q(s_t, a_t)$



Q-learning mouse example



Action values:

W	NW	N	NE	E	SE	S	SW
---	----	---	----	---	----	---	----


Choice:




Action values:

W	NW	N	NE	E	SE	S	SW
		+50					


Q-learning mouse example



Mouse cliff example



State:




Action values:

W	NW	N	NE	E	SE	S	SW
	-50	-50	-50	+50			

Mouse cliff example


- Q - learning results:



Mouse cliff example

What's wrong with Q-learning?



S-A-R-S-A




- Initialize $Q(s, a)$ with zeros
- Cycle:
 - Sample $\langle s, a, r, s', a' \rangle$ from environment
 - Compute $\hat{Q}(s, a) = r(s, a) + \gamma Q(s', a')$
 - Update: $Q(s_t, at) = \alpha \hat{Q}(s, a) + (1 - \alpha) Q(s_t, at)$

Mouse cliff example

- SARSA - learning results:





Expected value S-A-R-S-A



- Initialize $Q(s, a)$ with zeros
- Cycle:
 - Sample $\langle s, a, r, s' \rangle$ from environment
 - Compute $\hat{Q}(s, a) = r(s, a) + \gamma E_{a'' \sim \pi}[Q(s', a'')]$
 - Update: $Q(s_t, at) = \alpha \hat{Q}(s, a) + (1 - \alpha) Q(s_t, at)$

Mouse cliff example



S-A-R-S-A vs Q-learning

- SARSA gets optimal rewards under current policy
- Q-learning policy **would be** optimal (converges to Q^*)

S-A-R-S-A vs Q-learning

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

MC vs TD


- TD can learn *before* knowing the final outcome
 - TD can learn online after every step
 - MC must wait until end of episode before return is known
- TD can learn *without* the final outcome
 - TD can learn from incomplete sequences
 - MC can only learn from complete sequences
 - TD works in continuing (non-terminating) environments
 - MC only works for episodic (terminating) environments

MC vs TD: Bias/Variance Trade-off

- Return $G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$ is *unbiased* estimate of $v_\pi(S_t)$
- True TD target $R_{t+1} + \gamma v_\pi(S_{t+1})$ is *unbiased* estimate of $v_\pi(S_t)$
- TD target $R_{t+1} + \gamma V(S_{t+1})$ is *biased* estimate of $v_\pi(S_t)$
- TD target is much lower variance than the return:
 - Return depends on *many* random actions, transitions, rewards
 - TD target depends on *one* random action, transition, reward


Monte-Carlo Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$




Temporal-Difference Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$




Dynamic Programming Backup

$$V(S_t) \leftarrow \mathbb{E}_{\pi} [R_{t+1} + \gamma V(S_{t+1})]$$



Exploration/Exploitation Revisited

- Balance between using what you learned and trying to find something even better



Exploration/Exploitation Revisited

- Strategies:

- ϵ -greedy

- With probability ϵ take random action, otherwise take optimal action.

- Softmax

- Pick action proportional to softmax of shifted normalized Q-values.

$$\pi(a \mid s) = \text{softmax}(Q(s,a) / \tau)$$

- ϵ - dithering

- Adding random noise to Q-values with ϵ probability

Exploration/Exploitation over time

- If you want to converge to optimal policy you need to gradually reduce exploration.
- Example:
Initialize ϵ -greedy $\epsilon = 0.5$, then gradually reduce it
 - If $\epsilon \rightarrow 0$, it's **greedy in the limit**
 - Be careful with non-stationary environments

Resume

What have we learned today?

Questions?

